SRINIVAS UNIVERSITY

INSTITUTE OF ENGINEERING & TECHNOLOGY

Srinivas Campus,Mukka,Surathkal,Mangaluru-574 146

“Machine Learning and Ethical Challenges: A Case Study on Bias”

AN GROUP TASK

Submitted by:

NAME	**AASHITH.R.SHETTY**
USN	**01SU24CS004**

Subject:

Fundamentales of AI and ML

Subject code:

24SBT110

Submitted to:

Prof,Mahesh Kumar

Academic Year 2026

CONTENT

SSN	Section
1	Introduction
2	Case Study: Biased Hiring Tools
3	Analysis of Bias and Unintended Consequences
4	Proposed Solutions / Guidelines
5	Real-World Examples of ML Bias
6	Importance of Ethical AI Frameworks
7	Conclusion

Introduction

- Machine learning (ML) is increasingly used in **hiring, credit scoring, healthcare, law enforcement, and other decision-making systems**.
- While ML can improve efficiency, it can also **reinforce existing biases** if trained on historical or unrepresentative data.
- Ethical concerns arise when ML models **discriminate against individuals or groups** based on gender, race, age, or other sensitive attributes.
- Bias in ML can lead to **unfair outcomes, legal consequences, and loss of trust** in automated systems.
- Understanding and addressing bias is crucial for **responsible AI development**.
- ML ethics ensures that algorithms are **transparent, accountable, and equitable**.
- Proper data handling, fairness-aware modeling, and human oversight are essential to **minimize unintended consequences**.
- Applying ethical principles in ML helps **promote diversity, inclusion, and social responsibility** in AI systems.
- **Organizations must continuously audit AI systems** to detect and correct emerging biases.
- **Public awareness and regulatory guidelines** are important for ensuring ethical AI deployment.
- **Ethical ML practices contribute to sustainable, trustworthy AI** that benefits both individuals and society.

Case Study: Biased Hiring Tools

- Many companies use **ML-powered tools** to screen resumes and rank job candidates.
- **Problem:** Some hiring tools were trained on historical hiring data that favored **male candidates**, reflecting past human biases.
- As a result, **female candidates and other underrepresented groups** were unfairly penalized, even if qualified.
- The model **learned bias automatically** from historical patterns, without explicit instructions.
- Consequences included:
    - Reduced workforce **diversity**
    - Unfair **rejection of qualified candidates**

    - **Reputational damage** for companies using biased systems
- Publicized examples: Tech companies where AI models **discriminated against women in STEM roles**.
- Highlights the importance of **examining training data** and **feature selection** to prevent bias.
- Demonstrates that **automation does not eliminate human bias** if the data itself is biased.
- Ethical and legal issues can arise if biased ML models are **used at scale without oversight**.

Analysis of Bias and Unintended Consequences

- **Source of Bias:**
    - Historical hiring data contained **gender and demographic biases**.
    - Features such as previous job titles, schools, or resume wording **reflected past human preferences**.
- **Impact on Candidates:**
    - Qualified applicants from **underrepresented groups were unfairly scored lower**.
    - Reduced opportunities for women and minority candidates in certain roles.
- **Organizational Impact:**
    - Companies risk **legal challenges** due to discrimination claims.
    - Loss of **trust and reputation** among candidates and employees.
- **Propagation of Bias:**
    - If biased models are **reused or deployed at scale**, the problem compounds over time.
    - Bias can influence hiring patterns for **years to come**.
- **Unintended Consequences:**
    - Workforce becomes **less diverse**, limiting innovation and perspectives.
    - Ethical violations may go unnoticed if **ML decisions are not transparent**.
- **Key Takeaway:**
    - ML systems can **amplify existing social inequalities** if ethical safeguards are not applied.

Proposed Solutions / Guidelines

- **Data Auditing:**
  - Examine training data for **representation imbalances** and biased outcomes.
- **Feature Selection:**
  - Avoid using features that **encode sensitive attributes** such as gender, race, or age.
- **Bias Mitigation Techniques:**
  - Use **re-weighting or re-sampling** to balance the dataset.
  - Apply **fairness constraints** during model training to reduce bias.
- **Human Oversight:**
  - Combine ML recommendations with **human review** to reduce automated discrimination.
- **Transparency & Explainability:**
  - Ensure model decisions are **interpretable**, allowing stakeholders to understand outcomes.
- **Regular Monitoring:**
  - Continuously evaluate model performance for **emerging biases**.
- **Policy & Guidelines:**
  - Implement **ethical standards and regulatory compliance** for AI deployment in hiring.
- **Diverse Development Teams:**
  - Involve teams with **diverse backgrounds** to identify potential biases early.
- **Candidate Feedback Loops:**
  - Provide mechanisms for **applicants to appeal or review** decisions.
- **Documentation & Reporting:**
  - Maintain records of **data sources, model updates, and fairness audits** to ensure accountability.

1. Real-World Examples of ML Bias

- Showcase other cases where ML caused bias or unintended consequences.
    - **Credit Scoring:** Some AI systems denied loans disproportionately to certain racial groups.
    - **Facial Recognition:** Models misidentified individuals with darker skin tones more frequently.
    - **Healthcare Algorithms:** Certain predictive tools underestimated risk for minority patients.
- These examples emphasize the **widespread impact of biased AI**.

2. Importance of Ethical AI Frameworks

- Introduce **guiding principles** for ethical ML:
    - Fairness: Treat all groups equitably.
    - Accountability: Assign responsibility for model outcomes.
    - Transparency: Make models understandable to stakeholders.
    - Privacy: Protect sensitive data used in training.
- Highlights the need for **structured policies and governance** to prevent bias before deployment.

Conclusion

- Machine learning systems can **unintentionally reinforce biases** if trained on historical or unrepresentative data.
- Bias in ML has **real-world consequences**, affecting fairness, diversity, and trust in automated systems.
- Ethical AI requires **careful data handling, fairness-aware modeling, and human oversight**.
- Implementing guidelines such as **data auditing, transparency, and regular monitoring** can significantly reduce bias.

- Diverse development teams and continuous feedback mechanisms **help identify and mitigate unintended consequences**.
- Real-world examples, like biased hiring tools or unfair credit scoring, **highlight the importance of proactive measures**.
- Responsible ML ensures that AI systems are **transparent, accountable, equitable, and socially responsible**.
- Applying ethical frameworks in ML contributes to **sustainable, trustworthy AI** that benefits both individuals and society.